

Matrix multiplication: Strassen's algorithm

We've all learned the naive way to perform matrix multiplies in $O(n^3)$ time. In today's lecture, we review Strassen's sequential algorithm for matrix multiplication which requires $O(n^{\log_2 7}) = O(n^{2.81})$ operations;

Block Matrix Multiplication

The idea behind Strassen's algorithm is in the formulation of matrix multiplication as a recursive problem. We first cover a variant of the naive algorithm, formulated in terms of block matrices, and then parallelize it. Assume $A, B \in \mathbb{R}^{n \times n}$ and $C = AB$, where n is a power of two.

We write A and B as block matrices,

$$\begin{aligned} A &= \begin{pmatrix} A_{11} & A_{12} \\ A_{21} & A_{22} \end{pmatrix} \\ B &= \begin{pmatrix} B_{11} & B_{12} \\ B_{21} & B_{22} \end{pmatrix} \\ C &= \begin{pmatrix} C_{11} & C_{12} \\ C_{21} & C_{22} \end{pmatrix} \end{aligned}$$

where block matrices A_{ij} are of size $n/2 \times n/2$ (same with respect to block entries of B and C). Trivially, we may apply the definition of block-matrix multiplication to write down a formula for the block-entries of C , i.e.

$$C_{11} = A_{11}B_{11} + A_{12}B_{21}$$

$$C_{12} = A_{11}B_{12} + A_{12}B_{22}$$

$$C_{21} = A_{21}B_{11} + A_{22}B_{21}$$

$$C_{22} = A_{21}B_{12} + A_{22}B_{22}$$

Parallelizing the Algorithm

Realize that A_{ij} and B_{kl} are smaller matrices, hence we have broken down our initial problem of multiplying two $n \times n$ matrices into a problem requiring 8 matrix multiplies between matrices of size $n/2 \times n/2$, as well as a total of 4 matrix additions.

we may come up with the following recurrence for work:

$$T(n) = 8T(n/2) + O(n^2)$$

By the Master Theorem, $T(n) = O(n^{\log_2 8}) = O(n^3)$.

Strassen's Algorithm

We now turn toward Strassen's algorithm, such that we will be able to reduce the number of sub-calls to matrix-multiplies to 7, using just a bit of algebra. In this way, we bring the work down to $O(n^{\log_2 7})$.

We write down C_{ij} 's in terms of block matrices M_k 's. Each M_k may be calculated simply from products and sums of sub-blocks of A and B.

That is, we let

$$M1 = (A11 + A22) (B11 + B22)$$

$$M2 = (A21 + A22)B11$$

$$M3 = A11(B12 - B22)$$

$$M4 = A22(B21 - B11)$$

$$M5 = (A11 + A12)B22$$

$$M6 = (A21 - A11)(B11 + B12)$$

$$M7 = (A12 - A22)(B21 + B22)$$

It can be verified that

$$C11 = M1 + M4 - M5 + M7$$

$$C12 = M3 + M5$$

$$C21 = M2 + M4$$

$$C22 = M1 - M2 + M3 + M6$$

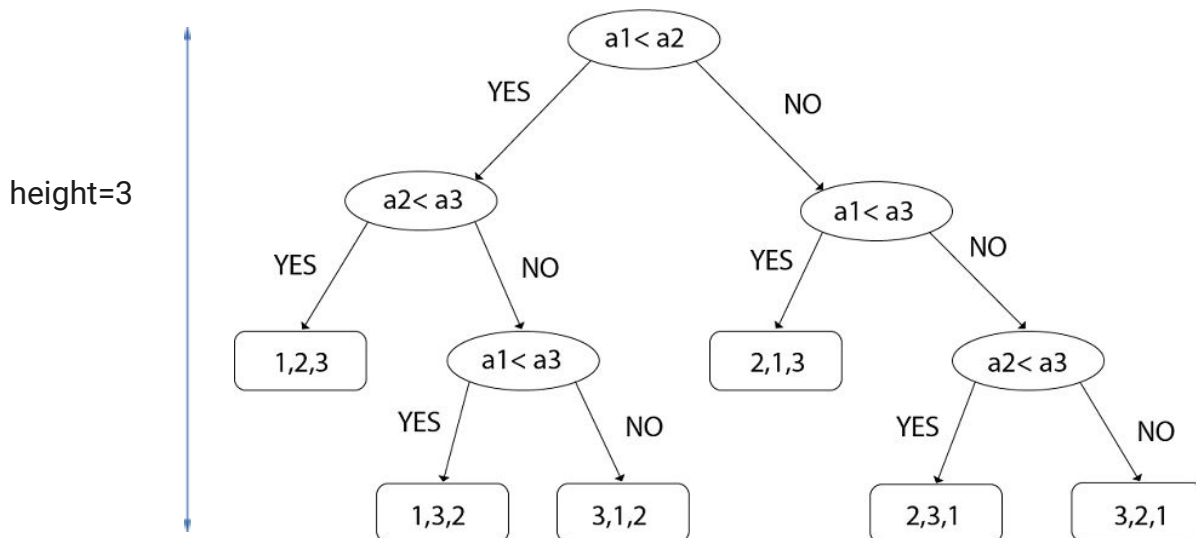
Realize that our algorithm requires quite a few summations, however, this number is a constant independent of the size of our matrix multiples. Hence, the work is given by a recurrence of the form

$$T(n) = 7T(n/2) + O(n^2) \Rightarrow T(n) = O(n^{\log_2 7}) = O(n^{2.81})$$

Lower Bound Theory for comparison sort

The preceding sections present three $O(n \log n)$ sorting algorithms—quick sort, heap sort, and the two-way merge sort. But is $O(n \log n)$ the best we can do?

In this section we answer the question by showing that any sorting algorithm that sorts using only binary comparisons must make $\Omega(n \log n)$ such comparisons.



Any sorting algorithm that uses only **binary comparisons** can be represented by a **binary decision tree**. Furthermore, it is the height of the binary decision tree that determines the worst-case running time of the algorithm.

Given an input sequence of **n items** to be sorted, every **binary decision tree** that correctly sorts the input sequence must have **at least $n!$** leaves—one for each permutation of the input. Therefore the **height of the binary decision tree** is at **least $\lceil \log_2 n! \rceil$** .

$$\begin{aligned}\lceil \log_2 n! \rceil &\geq \log_2 n! \\ &\geq \sum_{i=1}^n \log_2 i \\ &\geq \sum_{i=1}^{n/2} \log_2 i/2 \\ &\geq n/2 \log_2 n/2 \\ &= \Omega(n \log n)\end{aligned}$$